DTU Compute
Department of Applied Mathematics and Computer Science

Occupancy-Aware Climate Control Estimating Thermal Comfort Curves with Hidden Markov Models

Bachelor Thesis



Occupancy-Aware Climate Control
Estimating Thermal Comfort Curves with Hidden Markov Models

Bachelor Thesis
7 December, 2026

By
Mads Holme Håkonsson

Copyright: Reproduction of this publication in whole or in part must include the customary bibliographic citation, including author attribution, report title, etc.
Cover photo: Vibeke Hempler, 2012
Published by: DTU, Department of Applied Mathematics and Computer Science, Asmussens Alle, Building 322, 2800 Kgs. Lyngby Denmark
https://www.compute.dtu.dk
ISSN: [0000-0000] (electronic version)
ISBN: [000-00-0000-000-0] (electronic version)

ISSN: [0000-0000] (printed version)
ISBN: [000-00-0000-000-0] (printed version)

Approval

This thesis has been prepared over six months at the Section for Indoor Climate, Department of Civil Engineering, at the Technical University of Denmark, DTU, in partial fulfilment for the degree Master of Science in Engineering, MSc Eng.

It is assumed that the reader has a basic knowledge in the areas of statistics.

Mads Holme Håkonsson - s234936

.....

Signature

.....

Date

Abstract

ABSTRACT IS MISSING

Acknowledgements

Mads Holme Håkonsson, MSc Civil Engineering, DTU
Creator of this thesis template.

[Name], [Title], [affiliation]
[text]

[Name], [Title], [affiliation]
[text]

Contents

Preface ii
Abstract iii
Acknowledgements iv

1 Introduction **1**

2 Data **2**

3 Theory **6**
3.1 The Diffusion Process 6

Bibliography **9**

A Title **11**

1 Introduction

Reduction in the usage of energy in buildings is a crucial aspect if one wants to go towards a more sustainable future. It has long been known that buildings are responseible for X% of the total energy consumption in the world, and that a large part of this energy is for HVAC systems [MISSING REFERENCE] HVAC systems energy consumption has been directly related to the occupant behaviours of the buildings. This imposes a challenge, as one has to be able to understand the behaviour of the occupants in the buildings before we can gain the largest optimization in energy consumption. (MISSING REFERENCE) has shown that there is a large variability in the behavour of occupants such as thermal comfort, lighting & ventiliation needs and precense hours.

Within the context of energy consumption in buildings, heating has been estimated to contribute to X% [MISSING REFERENCE] of the consumption. An aspect of heating of a building is the byproduct of the thermal comfort of the occupants in the building, and therefore, if we can understand the thermal comforts of the occupants in the buildings, we can optimize the energy consumption drasitically. However, the thermal comfort curves is a personal preference of the occupants, and we should therefore expect this to differer dependent on the individual buildings and the number of occupants.

This projects aims to understand when the occupants of a single offices are present, and how much thermal comfort feedback does we need to collect from the occupants in order to be able to infer the true thermal comfort of the occupants in the building. The goal is to also build a codebase others and easyly download and extent.
We will do this by using a Hidden Markov Model (HMM) to model the state of the room

WRITTEN OUTLINE OF EACH CHAPTER AND WHAT QUESTIONS THEY ANSWER

2 Data

The data used in this project are CO_2 data collected from the office building 303B and the censors and data providers are climify [MISSING REFERENCE]. The data is collected in the fall and winter period of 2023, and the raw data period spans from [MISSING REFERENCE] to [MISSING REFERENCE]. The data collected from the censors are in intervals of around 10 minutes. Because our HMM models works in discrete time, the data have been aggregated into 30 minutes such that each time step $\Delta t = 1$ represents 30 minutes.

When collecting CO_2 data, the censors tend to drift over time [MISSING REFERENCE], and therefore, the data collected from the censors are not always accurate. To empty rooms, that have the same CO_2 levels may show different readins based on how much the censors have drifted over time. This is a problem, as we want to use the CO_2 data to infer the presence of the occupants in the building and a higher baseline of the CO_2 data may lead to a false positive of the presence of the occupants in the building. To correct for this, we have used a method called *Automatic Baseline Correction* (ABC) Algorithm. [MISSING REFERENCE]. In the use of this algorithm, we have assumed that the lowest CO_2 ppm values over a period is equal to the outside CO_2 ppm values. From literature, it's common to set the outside CO_2 ppm values to 400 ppm [MISSING REFERENCE], which we have also done.

Let C_t be the CO_2 ppm value at time t, $\tilde{C}_t$ be the corrected CO_2 ppm value at time t.

We then define the modified ABC algorithm as

$$\tilde{C}_t = C_t - Q_{0.02}\left(C_{t-\Delta t:t}\right) + 400, \tag{2.1}$$

where $Q_{0.02}\left(C_{t-\Delta t:t}\right)$ is the 2% quantile of the CO_2 ppm values over the time period $[t - \Delta t, t]$ and Δt is a fixed time period. In this project, we have set $\Delta t = 8$ days, such that the weekday overlaps. The reason for using a 2% quantile is to avoid the effect of erradic readings of the CO_2 ppm values, which may occur if the censors are not working properly. Because we are interrested in predicting future presence of the occupants in the building,

we dont need the data to be as granular as 10 minutes.

Room	N	Min	2.5%	Median	Mean	97.5%	Max
Room 001	14289	341.00	426.00	621.00	828.52	2411.00	3185.00
Room 003	14328	335.00	405.00	516.00	611.92	1424.00	3025.00
Room 004	14268	292.00	400.00	940.50	1201.19	3042.00	3710.00
Room 005	14324	345.00	418.00	704.00	774.84	1500.00	2447.00
Room 007	14334	360.00	414.00	633.00	736.45	1588.00	2805.00
Room 008	14321	360.00	429.00	781.00	988.87	2824.00	3581.00
Room 009	14346	325.00	417.00	685.00	935.13	2322.50	3654.00
Room 012	14344	335.00	414.00	739.00	769.75	1406.00	2860.00
Room 013	14363	247.00	442.00	716.00	852.76	1841.00	2579.00
Room 014	14144	327.00	422.00	804.00	1098.36	3305.42	4000.00
Room 015	14336	337.00	414.00	754.50	985.40	2706.25	3808.00
Room 016	14333	349.00	422.30	686.00	833.00	1841.70	3825.00
Room 019	14338	355.00	435.00	727.00	886.71	1927.57	2756.00

Table 2.1: Summary statistics of the raw CO_2 readings (ppm) per room, after removing missing, saturated and duplicate readings.

TO DO:

- Make summary table of the data with before and after the ABC algorithm (n, mean, sd, min, quantiles, max, % of bins below 450 ppm; one row per room if several rooms).
- Make ACF plots to show the AR structure of the data. Add PACF (shows AR order); note daily cycle at lag 48.
- Make a plot of the data before and after the ABC algorithm (raw, estimated floor, calibrated + zoom on one week).
- Histogram/density of calibrated CO_2 (multimodality motivates hidden states).
- Average daily profile, weekday vs. off-day (motivates off-day + Fourier covariates).
- Overview plot of full series with train/val/test boundaries and feedback events.
- Short chapter introduction and conclusion.

Fix in existing text:

- Δt used both as 30-min step and as 8-day window: rename window to W.
- Dates: Room 009 segment is 2023-11-21 to 2024-02-09 (3860 bins, $\approx$ 80 days), not "fall and winter 2023". Clarify 303B vs. DTU dataset (live `room_009/30min` data comes from `clean_dtu_data.py`).
- Reason for 8 days: every window contains a full weekend (an empty-room period).
- Merge the duplicated aggregation paragraph; state that bins are means of the readings.
- Typos: censors $\rightarrow$ sensors, "To empty rooms" $\rightarrow$ Two, readins, erradic, interrested, dont, climify $\rightarrow$ Climify.
- Define a `\COtwo` macro instead of retyping `$\text{CO}_2$`.

Missing processing steps to describe:

- Raw cleaning: drop NaNs, keep 0–4000 ppm (sensor saturation), remove duplicates.
- Gap handling: only the longest gap-free segment is kept (discrete model needs equal spacing).
- ABC details: applied after 30-min binning (and why); trailing window, so no leakage into val/test; first 8 days dropped as warm-up; justify 2% quantile (must be below share of empty bins, $\approx$16% below 500 ppm in Room 009).
- Report floor drift (Room 009: 388 $\rightarrow$ 500 ppm).
- State ABC assumption: room is empty and ventilated to outdoor level at some point in every window. Outdoor CO_2 was $\approx$420 ppm in 2023–24: justify 400 or mention it.
- Feedback data: how collected, number of responses, binary flag per 30-min bin.
- Humidity: calibration against outdoor humidity, or state that it is excluded.

- Covariates: off-day flag (weekends + holidays), daily (3 harmonics) and weekly (1 harmonic) Fourier terms, weather variables; z-scored with training statistics only.
- Time zone: timestamps in UTC (no DST change in Nov–Feb).
- Split: chronological 65/15/20, boundaries 2024-01-12 and 2024-01-24; Christmas break falls in training.
- Rooms and sensors: which rooms and why, single/shared offices, ventilation type, sensor type (NDIR), accuracy, sampling interval.

Missing references: Climify, NDIR sensor drift, ABC (Senseair firmware; check whether your sensors already run it), outdoor CO_2 baseline (e.g. NOAA).

3 Theory

In this chapter, a physical model of the underlying system will be presented and shown to be equivalent to a discrete-time AR(1)-HMM. A overview of the theory behind the models and inference algorithms will be presented as well.

3.1 The Diffusion Process

The underlying physical model of the CO_2 dynamics of a room can be described by a diffusion process. The diffusion process is a continuous-time stochastic process that describes the change of CO_2 in a room.
Let $\tilde{C(t)} = C(t) - \min_{\forall t}(C(t))$ such that the process reverts back to a baseline of 0. The diffusion process can then be desbribed as

$$\Delta\tilde{C(t)} = (P(t) - \alpha\tilde{C(t)})\Delta t + \sigma\Delta W(t), \tag{3.1}$$

where $P(t)$ is the CO_2 production rate of the occupants in the room, α is the decay rate of the CO_2 in the room, σ is the diffusion coefficient and $W(t)$ is a Wiener process, which is interpretated as the noise in the sensors.
This models tells us, that if the production rate of the occupants at time t is higher than the decay rate (or air change rate) of the room, then the CO_2 levels in the room will increase. The argument for this physical model that CO_2 accumate and that it does not decay by itself [MISSING REFERENCE], unless it's ventilated out of the room.

3.1.1 Proof of AR(1) process

We now want to show, that the underlying physical model of the CO_2 dynamics of a room can be described as a discrete-time AR(1) process.
Firstly, we recognize the diffusion process as a Ornstein-Uhlenbeck, and the solution of this process is given by Iacus (2008, Sec. 1.13.1) as

$$\tilde{C}(t) = \frac{P}{\alpha} + \left(x_0 - \frac{P}{\alpha}\right)e^{-\alpha t} + \theta_3 \int_0^t e^{-\alpha(t-u)}\,\mathrm{d}W_u \tag{3.2}$$

or, in the second parametrization,

$$\tilde{C}(t) = \mu + (x_0 - \mu)e^{-\alpha t} + \sigma \int_0^t e^{-\alpha(t-u)}\,\mathrm{d}W_u, \qquad \mu = \frac{P}{\alpha}. \tag{3.3}$$

where x_0 is the initial value of the process at time $t = 0$.
If we then discretize the process with a time step Δt, we can write the process as

$$\tilde{C}_{t+\Delta t} = \mu + (\tilde{C}_0 - \mu)e^{-\alpha(t+\Delta t)} + \sigma \int_0^{t+\Delta t} e^{-\alpha(t+\Delta t-u)}\,\mathrm{d} \tag{3.4a}$$

where we can split the interal into two parts as

$$\sigma \int_0^{t+\Delta t} e^{-\alpha(t+\Delta t-u)}\,\mathrm{d}W_u = \sigma \left(e^{-\alpha\Delta t} \int_0^t e^{-\alpha(t-u)}\,\mathrm{d}W_u + \right. \tag{3.4b}$$

We can then write out the process as

$$\tilde{C}_{t+\Delta t} = (1 - e^{-}\mu + (\tilde{C}_0 - \mu)e^{-\alpha t}e^{-\alpha \Delta t}$$

$$+\sigma e^{-\alpha \Delta t} \int_0^t e^{-\alpha(t-u)}\, \mathrm{d}W_u + \sigma \int_t^{t+\Delta t} e^{-\alpha(t+\Delta t-u)}\, \mathrm{d}W_u \tag{3.4c}$$

due to $\delta t > 0$ and $\alpha > 0$

(3.4d)

3.1.2 Modelling Production Rate

By assuming that the prodction rate of the occupants is only depented on the amount of occupants in the room, we can model the production rate as a function of the occupancy state of the room. Let $S(t)$ be the occupancy state of the room at time t, then we can model the production rate as

$$P(t) = \beta_{S(t)}, \tag{3.5}$$

where $\beta_{S(t)}$ is the production rate of the occupants in the room, which is dependent on the occupancy state of the room.

Bibliography

Iacus, Stefano M. (2008). *Simulation and Inference for Stochastic Differential Equations: With R Examples*. Springer Series in Statistics. New York: Springer. ISBN: 978-0-387-75838-1. DOI: 10.1007/978-0-387-75839-8.

A Title

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Technical
University of
Denmark

Asmussens Alle, Building 322
2800 Kgs. Lyngby
Tlf. 4525 1700

https://www.compute.dtu.dk